

Now that you know the fundamentals of back prop...

## Advanced Back Propagation

“Deep” Neural Nets

“Convolutional” Neural Nets

Recurrent Neural Nets

1

## The Autoencoder

Output nodes:



$\mathcal{N}$  of these

Hidden nodes:



$< \mathcal{N}$  of these (typically)


Input nodes:



$\mathcal{N}$  of these


2

## The Autoencoder




3

## The Autoencoder: Training




4

## The Autoencoder: Training




5

## The Autoencoder: Training




6

## The Autoencoder: Training




7

## The Autoencoder: Result



8

**The Point:** An Autoencoder discovers statistical regularities in its inputs and encodes them in its “hidden” layer...



Now, let's change gears for a moment...

9

## What's Wrong With Back Prop?

(a gazillion things. But if you're an engineer...)

- The  $\Delta w_{ij} = f(\text{error} * a_i * (1 - a_i))$  derivative:
  - Goes to *zero* as activation approaches zero or one.
- Happens in each layer *successively*
- Result: *Errors get washed out to near nothing after just a few layers.*
  - *Many-layered perceptrons therefore take forever to train.*
- But some problems are solved more easily with more layers ☹
  - What is a poor ~~hack~~ computer programmer to do??
  - This problem made back prop ~~unprofitable~~ unpopular for several years. *Until...*

10


## An Idea!

- Geoff Hinton: *What if we pre-train a neural net with an auto-encoder, and then use **that** as the input layer to another network?*
  - (This is the idea you had earlier)
- *And what if we do this over and over?!*
- *Like, with lots of hidden nodes and lots of hidden layers?*
- *Kinda like...*

11


## An Idea!

Start with an autoencoder, *but...*



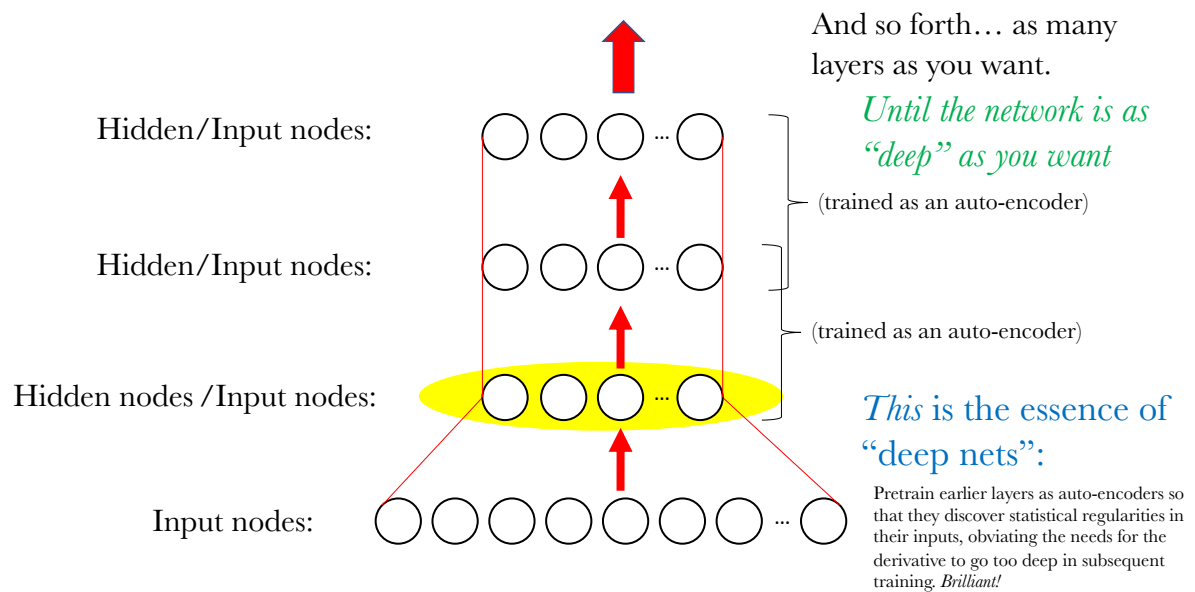
12

## An Idea!



13

## An Idea!



14

## “Convolutional” Nets

It’s just weight duplication (aka “weight sharing”, “error sharing”, “update sharing”).

Start with an image.

Take, say, a 10 X 10 pixel piece of the image...


and train an auto-encoder on it.



15

### At each 10 X 10 Window Over the Image...


Output Layer:



Hidden Layer:



Input Layer:



16



## “Convolutional” Nets

It’s just weight duplication (aka “weight sharing”, “error sharing”, “update sharing”).

Start with an image.

Take, say, a 10 X 10 pixel piece of the image...

and train an auto-encoder on it.

Move it over 5 pixels and repeat. Over and over. Until you’ve covered the whole image.



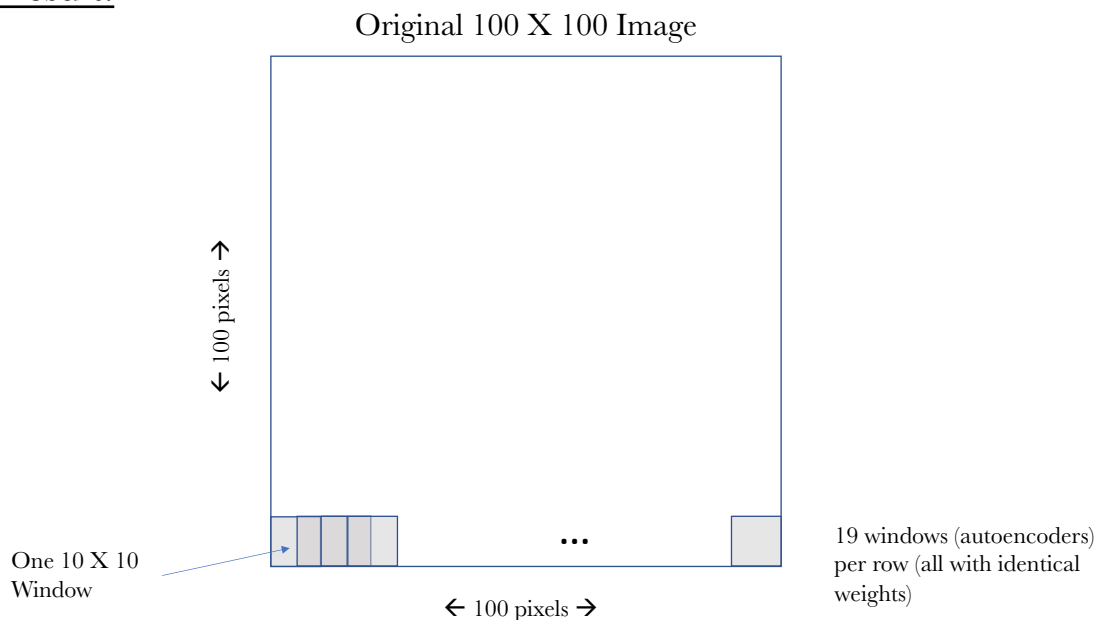
Repeat for multiple images.

The trick: *Weight sharing*: It’s the same weight matrix over and over. One weight matrix gets to learn from every location.

Result (sometimes): The weight matrix learns “V1-like” RFs.

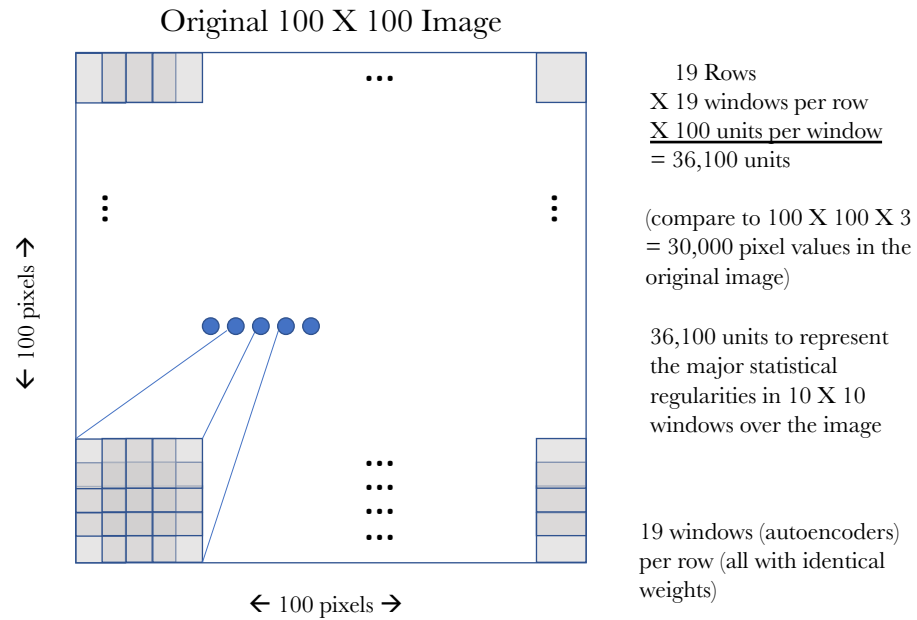
17

### The Result:



18

## The Result:



19

## “Convolutional” Nets


Armed with that recoding of the image (an  $N \times N$  matrix of  $M$  “V1-like” units)...

- Use *that* as the input to a “deep” net...
- Or do “pooling” over adjacent locations (*not* Poggio’s idea; see Fukushima & Miyaki, 1982)...
- Or train each location with an autoencoder...
- Or train multiple adjacent locations with an autoencoder...
- Or do... whatever your little capitalist heart desires.
- *Much* money to be made here, but no intellectual progress. (An exercise for dupes. Preferably dupes with money.)

20

## Recurrent Nets

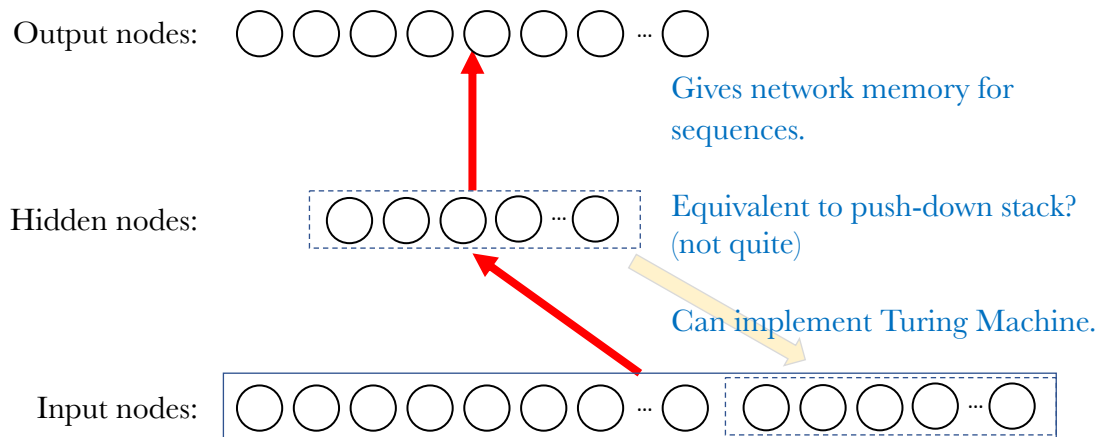
An ordinary back prop net, *but...*



21

## Recurrent Nets

An ordinary back prop net, *but...*



22